

BRANDON MINER

San Francisco, CA

brandonm333@outlook.com +1 408-335-5436 [linkedin.com/in/brandon-miner-3x3](https://www.linkedin.com/in/brandon-miner-3x3) github.com/Brandonminer333/

Education

University of San Francisco

M.S. in Data Science and Artificial Intelligence

San Francisco, CA

Jun 2026

San Diego State University

B.S. in Statistics with Emphasis in Data Science, Minor in Mathematics

San Diego, CA

Dec 2025

Experience

Data Scientist

San Francisco, CA

ACLU *Oct 2025 – Present*

- Architecting the software design of an end-to-end **multi-modal AI** video auditing platform from scratch (**Python, PyTorch, FastAPI, MySQL**) to automate analysis of unstructured police bodycam footage, projected to eliminate **300+ hours** of manual review per monthly cycle, all while tracking product development.
- Engineered a **GPU-accelerated inference pipeline** using **FFmpeg, SpeechBrain metricgan-plus-voicebank, Pyannote speaker-diarization-3.1**, and **OpenAI Whisper** for specialized speech-to-text transcription, processing **600+ videos per month**
- Organized a suite of **multi-stage NLP classifiers and extractors** utilizing **LLM-driven auto-labeling** to generate training data for fine-tuning to extract audit signals from transcripts at scale.
- Designed and deployed a **private-server FastAPI backend** paired with a **React.js** frontend, delivering the full AI auditing system as a production-ready web application for non-technical end-users.

Data Scientist, Lead Intern

San Diego, CA

San Diego County Taxpayers Association

Mar 2024 – Aug 2025

- Led data science initiatives, managing a team of **4+ associates** to run statistical data analysis to deliver actionable ML-driven insights on civil workforce and homelessness programs, utilizing leadership and communication skills.
- Built a workforce analytics pipeline combining iterative **web scraping** and **NLP-based entity standardization** to analyze wage dynamics across public records, surfacing a statistically significant **19.7%** police wage increase per **1%** staffing drop.
- Applied **mixed-effects regression models** and **multicollinearity reduction** to analyze regional homelessness expenditure data, identifying high-impact interventions (e.g., rapid rehousing) to guide future resource allocation decisions.

Technical Projects

Parkcast-SF: San Francisco Parking Assistant

- Designed the end-to-end architecture for a San Francisco parking assistant app with a Vercel frontend and Google Cloud Run API backend.
- Predicted parking availability from historical meter data and used **LightGBM** to predict metered parking by block; paired with **K-nearest neighbors** to approximate availability in non-metered zones.
- Implemented a challenger/champion model lifecycle with MLflow-backed evaluation and new challengers deployed every **3 days** to improve production performance.
- Simulated mapping UX using OpenStreetMap and Mapbox to support location-based navigation without Google Maps dependency.

California Grape ETL Pipeline & Cloud Dashboard

- Engineered a multi-source ETL pipeline** in Python to extract, transform, and consolidate disparate agricultural datasets into a unified schema for downstream ML and analytics.
- Containerized the full-stack application** using **Docker Compose** and deployed a serverless instance to **Google Cloud Run**, ensuring scalable, reliable cloud hosting.

Technical Skills

Programming Languages: Python (Pandas, NumPy, Scikit-learn, TensorFlow, PyTorch, FastAPI, Transformers, Hugging Face, Streamlit), SQL (PostgreSQL), R (tidyverse), Spark SQL

AI Systems: LLM Integration & Prompting, LLM-driven Auto-labeling, Multi-modal AI, OpenAI Whisper, FFmpeg, GPU-Accelerated Inference, Docker, Google Cloud Platform, Git

Machine Learning: Deep Learning (Neural Networks, Transformers, Fine-tuning), NLP (ASR, Text Classification), Regression, Random Forests, Bayesian Text Classification, PCA, Explainable AI (XAI), Distributed Computing (Spark)