

BRANDON MINER

San Francisco, CA

brandonm333@outlook.com linkedin.com/in/brandon-miner-3x3 github.com/Brandonminer333/

Forward Deployed Engineer — End-to-end ML systems on customer infrastructure — from stakeholder needs through production deployment

Education

University of San Francisco

M.S. in Data Science and Artificial Intelligence

San Francisco, CA

Expected Jun 2026

- Distributed Computing
- Machine Learning
- Communication
- Mlops

San Diego State University

B.S. in Statistics with Emphasis in Data Science, Minor in Mathematics

San Diego, CA

Dec 2025

Experience

Forward Deployed Engineer

ACLU

San Francisco, CA

Oct 2025 – Present

- Owned end-to-end delivery of a production **ML inference platform** (Python, FastAPI, MySQL, Docker) deployed directly on client hardware — from requirements gathering through deployment — eliminating **300+ hours** of manual review per monthly cycle.
- Embedded with non-technical stakeholders to translate operational needs into a **GPU-accelerated ETL and inference pipeline** using **FFmpeg** and **Whisper ASR**, processing **600+ videos per month** with zero cloud dependency.
- Architected and shipped a **multi-stage NLP classification system** combining LLM-assisted auto-labeling, iterative fine-tuning, and structured output parsing — delivering audit-ready signals directly in the hands of end-users.
- Deployed the full ML stack on a **private-server FastAPI backend** with a **React.js** web frontend, enabling non-technical civil rights investigators to run inference without any engineering support.

Data Scientist, Lead Intern

San Diego County Taxpayers Association

San Diego, CA

Mar 2024 – Aug 2025

- Led a cross-functional team of **4 analysts** to build data pipelines and deliver analytical products directly to government stakeholders, translating messy public datasets into actionable policy insights.
- Designed and deployed an end-to-end **data standardization pipeline** combining iterative web scraping and text normalization; applied regression modeling to quantify a **19.7% police wage increase** per **1% staffing reduction** — findings presented to county officials.
- Applied **mixed-effects regression** with VIF-based multicollinearity reduction to regional homelessness expenditure data; surfaced high-impact interventions to guide resource allocation decisions for program leadership.

Technical Projects

GenAI Franchise Personality Quiz

- Built an end-to-end quiz generator: users describe any fictional class system in natural language; Gemini parses the franchise, scrapes Fandom wikis for canon characters, and roleplays each through 15 Likert questions to produce reference vectors asynchronously via FastAPI background jobs.
- Deployed a split architecture on Vercel (frontend) and Google Cloud Run (API), with durable quiz artifacts in GCS, shareable quiz URLs, a community quiz library, and GitHub Actions CI (pre-commit, pytest, automated API deploy).

California Grape ETL Pipeline & Dashboard

- Built a **multi-source ETL pipeline** in Python to extract, transform, and consolidate disparate agricultural datasets into a unified schema, demonstrating ability to rapidly onboard domain-specific data and deliver working systems.
- **Containerized the full-stack application** with **Docker Compose** and deployed to **Google Cloud Run** as a scalable serverless instance — reproducible infrastructure ready for client-site or cloud deployment.
- Delivered an interactive **Streamlit dashboard** for dynamic visualization of regional production metrics, providing a user-facing interface accessible to non-technical stakeholders.

Parkcast-SF: San Francisco Parking Assistant

- Owned system design and the production **FastAPI** inference API for a full-stack USF capstone parking platform, serving block-level **LightGBM** occupancy predictions from artifacts in **GCS**.
- Applied **K-nearest neighbors**-derived baselines to support **LightGBM** inference on sparse, non-metered blocks where direct sensor history was unavailable.
- Built automated model **CI/CD** on **GCP**: scheduled **GitHub Actions** retraining with **MLflow** experiment tracking and a **champion/challenger** promotion gate (validation MAE) before GCS upload; split deploy across **Vercel** (frontend) and **Cloud Run** (API).

Technical Skills

Programming & Frameworks: Python (FastAPI, PyTorch, Pandas), JavaScript (React.js), SQL (PostgreSQL), Bash

Delivery & Infrastructure: Docker, Docker Compose, Google Cloud Run, private-server deployment, client-hardware deployment, GPU-accelerated inference, REST APIs, GitHub Actions CI/CD

ML Systems: LLM-assisted labeling, NLP pipelines, multi-source ETL, batch inference, stakeholder requirements gathering

Soft Skills: customer embedding, cross-functional leadership, communication with non-technical stakeholders