

BRANDON MINER

San Francisco, CA

brandonm333@outlook.com linkedin.com/in/brandon-miner-3x3 github.com/Brandonminer333/

Data Scientist \$—\$ Predictive analytics, anomaly detection, and reproducible ML models on large structured and unstructured datasets

Education

University of San Francisco

M.S. in Data Science and Artificial Intelligence

San Francisco, CA

Jul 2025 – Jun 2026

San Diego State University

B.S. in Statistics with Emphasis in Data Science, Minor in Mathematics

San Diego, CA

Aug 2022 – Dec 2025

Experience

Data Scientist

ACLU

San Francisco, CA

Oct 2025 – Jun 2026

- Architecting the product development of an end-to-end ML video auditing platform from scratch (**Python, FastAPI, MySQL, Docker**) to automate data processing and analysis of unstructured police bodycam footage, projected to eliminate **300+ hours** of manual review per monthly cycle.
- Developing a GPU-accelerated **ETL pipeline** using **FFmpeg** and **Whisper**-based speech-to-text to ingest, transcribe, and process batches of **600+ videos per month**; integrated **LLM-driven auto-labeling** to generate fine-tuning data for multi-stage text classification models extracting actionable audit signals.
- Designed and trained **multi-stage NLP classifiers** on auto-labeled transcripts to detect use-of-force indicators, rights violations, and officer conduct patterns at scale.
- Built a **private-server FastAPI backend** and **React.js** web frontend to deliver the full auditing system to non-technical stakeholders, enabling self-service model inference without engineering overhead.

Data Scientist, Lead Intern

San Diego County Taxpayers Association

San Diego, CA

Mar 2024 – Aug 2025

- Led data science initiatives, managing a team of **4+ associates** to deliver statistical insights on civil workforce economics and regional homelessness programs using statistical data analysis.
- Built a public workforce analytics pipeline combining **iterative web scraping** and **NLP-based entity standardization** to clean and link heterogeneous salary records; applied regression modeling to quantify a **19.7% police wage increase per 1% staffing drop**.
- Applied **mixed-effects regression models** with multicollinearity reduction (VIF analysis) to regional homelessness expenditure data, surfacing high-impact interventions such as rapid rehousing to guide resource allocation decisions.

Technical Projects

Credit Card Fraud Detection

- Built a binary classification model using **Logistic Regression (Scikit-Learn)** to identify fraudulent transactions on a highly imbalanced dataset.
- Tuned classification threshold via **ROC-AUC analysis** to minimize False Negative Rate (missed fraud) while capping False Positive Rate at **under 2%**.

Parkcast-SF: San Francisco Parking Assistant

- Built automated model **CI/CD** on **GCP**: scheduled **GitHub Actions** retraining with **MLflow** experiment tracking and a **champion/challenger** promotion gate (validation MAE) before GCS upload; split deploy across **Vercel** (frontend) and **Cloud Run** (API).
- Predicted parking availability across SF using historical parking meter records; where meters were unavailable, applied a **K-nearest neighbors** approximation to infer nearby parking availability.

Technical Skills

Programming & Tools: Python (Pandas, NumPy, Scikit-learn, PyTorch, LightGBM, Streamlit), R (tidyverse), SQL (PostgreSQL), Bash

Modeling & Analytics: predictive modeling, anomaly and failure-signal detection, classification, regression, mixed-effects models, NLP pipelines, PCA, ROC-AUC threshold optimization, imbalanced classification, experimental design

Data Engineering: large-scale ETL pipelines, unstructured text processing (ASR, NLP), multi-source data integration, reproducible model workflows (MLflow)

Communication & Leadership: technical documentation, stakeholder-facing analytics, cross-functional collaboration, team leadership