Brandon Miner

San Francisco, CA
brandonm333@outlook.com linkedin.com/in/brandon-miner-3x3 github.com/Brandonminer333/

Data Scientist — Statistical modeling, NLP pipelines, and end-to-end ML products in civic and public-interest contexts

Education

University of San Francisco — **San Francisco, CA**
M.S. in Data Science and Artificial Intelligence — *Expected Jun 2026*

- *Distributed Computing*
- *Machine Learning*
- *Deep Learning*
- *Experiments* / Experiments in Data Science

San Diego State University — **San Diego, CA**
B.S. in Statistics with Emphasis in Data Science, Minor in Mathematics — *Dec 2025*

Experience

Data Scientist — **San Francisco, CA**
ACLU — *Oct 2025 – Present*

- Architecting the product development of an end-to-end ML video auditing platform from scratch (**Python, PyTorch, FastAPI, PostgreSQL, Docker**) to automate data processing and analysis of unstructured police bodycam footage, projected to eliminate **300+ hours** of manual review per monthly cycle.
- Developing a GPU-accelerated **ETL pipeline** using **FFmpeg** and **Whisper**-based speech-to-text to ingest, transcribe, and process batches of **600+ videos per month**; integrated **LLM-driven auto-labeling** to generate fine-tuning data for multi-stage text classification models extracting actionable audit signals.
- Designed and trained **multi-stage NLP classifiers** on auto-labeled transcripts to detect use-of-force indicators, rights violations, and officer conduct patterns at scale.
- Built a **private-server FastAPI backend** and **React.js** web frontend to deliver the full auditing system to non-technical stakeholders, enabling self-service model inference without engineering overhead.

Data Scientist, Lead Intern — **San Diego, CA**
San Diego County Taxpayers Association — *Mar 2024 – Aug 2025*

- Led data science initiatives, managing a team of **4+ associates** to deliver statistical insights on civil workforce economics and regional homelessness programs using statistical data analysis.
- Built a public workforce analytics pipeline combining **iterative web scraping** and **NLP-based entity standardization** to clean and link heterogeneous salary records; applied regression modeling to quantify a **19.7% police wage increase per 1% staffing drop**.
- Applied **mixed-effects regression models** with multicollinearity reduction (VIF analysis) to regional homelessness expenditure data, surfacing high-impact interventions such as rapid rehousing to guide resource allocation decisions.

Technical Projects

California Grape ETL Pipeline & Dashboard

- Engineered a **multi-source ETL pipeline** in Python to extract, transform, and consolidate disparate agricultural datasets into a unified schema optimized for downstream analysis.
- **Developed an interactive Streamlit front-end** dashboard for real-time filtering and visualization of regional production metrics.
- Developed an **interactive Streamlit dashboard** enabling end-users to dynamically visualize and filter regional production metrics; view live deployment.

Credit Card Fraud Detection

- Built a binary classification model using **Logistic Regression (Scikit-Learn)** to identify fraudulent transactions on a highly imbalanced dataset.
- Tuned classification threshold via **ROC-AUC analysis** to minimize False Negative Rate (missed fraud) while capping False Positive Rate at **under 2%**.

Technical Skills

Programming & Statistical Tools: Python (Pandas, NumPy, Scikit-learn, PyTorch, Streamlit), R (tidyverse), SQL (PostgreSQL), Bash
Modeling & Analysis: regression, mixed-effects models, NLP pipelines, classification, PCA, multicollinearity reduction (VIF), ROC-AUC threshold optimization, imbalanced classification, experimental design
Data Engineering: multi-source ETL pipelines, web scraping, text normalization, entity standardization
Communication & Delivery: Streamlit dashboards, stakeholder-facing analytics, team leadership