

BRANDON MINER

San Francisco, CA

brandonm333@outlook.com +1 408-335-5436 [linkedin.com/in/brandon-miner-3x3](https://www.linkedin.com/in/brandon-miner-3x3) github.com/Brandonminer333/

Forward Deployed Engineer | End-to-end ML systems on customer infrastructure — from stakeholder needs through production deployment

Education

University of San Francisco

San Francisco, CA

M.S. in Data Science and Artificial Intelligence

Expected Jun 2026

- *Distributed Computing*
- *Machine Learning*
- *Communication*
- *Mlops*

San Diego State University

San Diego, CA

B.S. in Statistics with Emphasis in Data Science, Minor in Mathematics

Dec 2025

Experience

Forward Deployed Engineer

San Francisco, CA

ACLU

Oct 2025 – Present

- Owned end-to-end delivery of a production **ML inference platform** (Python, PyTorch, FastAPI, PostgreSQL, Docker) deployed directly on client hardware — from requirements gathering through deployment — eliminating **300+ hours** of manual review per monthly cycle.
- Embedded with non-technical stakeholders to translate operational needs into a **GPU-accelerated ETL and inference pipeline** using **FFmpeg** and **Whisper ASR**, processing **600+ videos per month** with zero cloud dependency.
- Architected and shipped a **multi-stage NLP classification system** combining LLM-assisted auto-labeling, iterative fine-tuning, and structured output parsing — delivering audit-ready signals directly in the hands of end-users.
- Deployed the full ML stack on a **private-server FastAPI backend** with a **React.js** web frontend, enabling non-technical civil rights investigators to run inference without any engineering support.

Data Scientist, Lead Intern

San Diego, CA

San Diego County Taxpayers Association

Mar 2024 – Aug 2025

- Led a cross-functional team of **4 analysts** to build data pipelines and deliver analytical products directly to government stakeholders, translating messy public datasets into actionable policy insights.
- Designed and deployed an end-to-end **data standardization pipeline** combining iterative web scraping and text normalization; applied regression modeling to quantify a **19.7% police wage increase** per **1% staffing reduction** — findings presented to county officials.
- Applied **mixed-effects regression** with VIF-based multicollinearity reduction to regional homelessness expenditure data; surfaced high-impact interventions to guide resource allocation decisions for program leadership.

Technical Projects

GenAI Franchise Personality Quiz

- Built an end-to-end quiz generator: users describe any fictional class system in natural language; Gemini parses the franchise, scrapes Fandom wikis for canon characters, and roleplays each through 15 Likert questions to produce reference vectors asynchronously via FastAPI background jobs.
- Deployed a split architecture on Vercel (frontend) and Google Cloud Run (API), with durable quiz artifacts in GCS, shareable quiz URLs, a community quiz library, and GitHub Actions CI (pre-commit, pytest, automated API deploy).

California Grape ETL Pipeline & Dashboard

- Built a **multi-source ETL pipeline** in Python to extract, transform, and consolidate disparate agricultural datasets into a unified schema, demonstrating ability to rapidly onboard domain-specific data and deliver working systems.
- **Containerized the full-stack application** with **Docker Compose** and deployed to **Google Cloud Run** as a scalable serverless instance — reproducible infrastructure ready for client-site or cloud deployment.
- Delivered an interactive **Streamlit dashboard** for dynamic visualization of regional production metrics, providing a user-facing interface accessible to non-technical stakeholders.

Parkcast-SF: San Francisco Parking Assistant

- Owned system design and the production **FastAPI** inference API for a full-stack USF capstone parking platform, serving block-level **LightGBM** occupancy predictions from artifacts in **GCS**.
- Applied **K-nearest neighbors**-derived baselines to support **LightGBM** inference on sparse, non-metered blocks where direct sensor history was unavailable.
- Built automated model **CI/CD** on **GCP**: scheduled **GitHub Actions** retraining with **MLflow** experiment tracking and a **champion/challenger** promotion gate (validation MAE) before GCS upload; split deploy across **Vercel** (frontend) and **Cloud Run** (API).

Technical Skills

Programming & Frameworks: Python (FastAPI, PyTorch, Pandas), JavaScript (React.js), SQL (PostgreSQL), Bash

Delivery & Infrastructure: Docker, Docker Compose, Google Cloud Run, private-server deployment, client-hardware deployment, GPU-accelerated inference, REST APIs, GitHub Actions CI/CD

ML Systems: LLM-assisted labeling, NLP pipelines, multi-source ETL, batch inference, stakeholder requirements gathering

Soft Skills: customer embedding, cross-functional leadership, communication with non-technical stakeholders